


Recap of Operating Systems

- The following OS concepts are required to understand virtualization
 - The concept of a process
 - Virtual memory, paging, (segmentation)
 - User mode and kernel mode of a process
 - Interrupt/trap processing in kernel mode
 - I/O handling

The concept of a process

- A process is a **running program**
- User writes a program, compiles it to generate an **executable**
- Executable contains **machine/CPU instructions**
 - Every CPU architecture (e.g., x86) defines a certain set of instructions
 - Compiler translates high level language code to instructions the CPU can run
- To create a process, OS allocates memory in RAM for the **memory image of the process**, containing
 - **Code** and static/global **data** from the executable
 - **Heap** memory for dynamic memory allocations (e.g., malloc)
 - **Stack** to store arguments/return address/local variables during function calls
- All instructions and variables in the memory image are assigned **memory addresses**
 - Starting at 0, up to some max value (4GB in 32-bit systems)

Memory image of a process



from compiled executable
(machine instructions, static/global data)

Address of this memory returned by malloc


↔ Heap and stack grow towards each other

New stack frame pushed by function call
Contains arguments, local vars, return address etc.
Popped when function returns

Process execution

- When a process is run, CPU executes the code in the memory image
- When a process runs on the CPU, the **CPU registers** hold values related to the process execution
 - The **program counter** (PC, or EIP in x86) has address of current instruction
 - The CPU fetches the current instruction, decodes it, and executes it
 - Any variables needed for operations are loaded from process memory into **general purpose CPU registers** (EAX, EBX, ECX, EDX etc in x86)
 - After instruction completes, values are stored from registers into memory
 - The **stack pointer** (SP, or ESP in x86) has address of top of stack (current stack frame holds arguments/variables of the current function that is running)
- The set of values of all CPU registers pertaining to a process execution is called its **CPU context**


CPU context during execution





Concurrent execution, context switching

- To run a process, OS allocates memory, loads CPU context
 - EIP points to instructions, ESP points to stack of process, registers have process data
 - CPU now begins to run the process
- OS runs multiple process **concurrently** by multiplexing on the same CPU
- **Context switch:** After running a process for some time, OS switches from one process to another
- How does context switch happen? OS saves the CPU context of the old process and loads the context of new process
 - When EIP points to instruction of new process, new process starts to run
- Where is the context saved?
 - OS has a data structure called **Process Control Block (PCB)** for each process
 - PCB (specifically, a PCB field called kernel stack) temporarily stores context of a process when it is not running

Context switch





CPU has switched execution from blue process to green process



User memory

Kernel memory


Process control blocks



Virtual memory

- Addresses assigned in memory image, that are used by CPU to load and store are **virtual/logical addresses**
 - These virtual addresses are not actual addresses occupied by instructions/data of process, but assigned from 0 for convenience
- Actual addresses where instructions/data bytes of process are stored are called **physical addresses**
 - RAM hardware needs physical addresses to fetch bytes
- CPU requests code/data at virtual addresses
 - Translated to physical address so that RAM can fetch it
- Memory is allocated at granularity of **pages** (usually 4KB)
 - Logical pages of a process are stored in physical frames in memory
 - Logical page numbers translated to physical frame numbers


Virtual and physical address spaces



Address translation with paging


- On every memory access, a piece of hardware called **MMU (memory management unit)** translates virtual addresses to physical addresses
- Page table of a process stores the **mapping of logical page number → physical frame number**
 - OS builds the page table when allocation memory
 - MMU uses this page table to translate addresses
- MMU looks up **CR3 register of CPU (x86)** which stores location of page table of current process
 - Looks up page number in page table to translate address
 - CR3 reset on every context switch
- Recent address translations cached in **TLB** (Translation Lookaside Buffer) located within MMU

Page table
(page 1 = frame y)



Virtual address space

Physical address space



Fetch address 5KB

Fetch address Y+1KB

Address = X

Address = Y

Address = Z

Virtual address



Using page table


Physical address




Hierarchical page tables

- Modern operating systems use hierarchical page tables
- 32-bit virtual address space (4GB), 2^{12} byte (4KB) pages
 - Each process can have up to $2^{32}/2^{12} = 2^{20}$ pages
 - Each page has a page table entry (PTE), so 2^{20} PTEs per process
 - Assuming each PTE is 4 bytes, all page table entries occupy $4 * 2^{20} = 4\text{MB}$
- Cannot store a large page table contiguously in memory, so page table of a process is stored in memory in page sized chunks
 - Each 4KB page stores 2^{10} PTEs, so $2^{20}/2^{10} = 2^{10}$ pages to store all PTEs
 - Pointers to these “inner” pages stored in an outer page directory
 - In 32-bit architectures, one outer page can store physical frame numbers of all 2^{10} inner page table pages, so 2-level page table
 - More levels in page table for 64-bit architectures

Page table lookup (walking the page table)



Demand Paging




User mode and kernel mode

- Two kinds of CPU instructions
 - Unprivileged instructions (regular program code) – part of user code
 - **Privileged instructions** (access to hardware etc.) – part of OS code
- Modern CPUs have multiple privilege levels (**rings**)
 - Privileged instructions are executed only when CPU is at high privilege level
- User code runs at low privilege level / **user mode** (CPU set to ring 3 in x86)
 - If user runs privileged instructions, error is thrown
- OS code runs at high privilege level / **kernel mode** (CPU set to ring 0 in x86)
 - Allowed to run privileged instructions
- Privilege level checked by CPU and MMU (every page has privilege bit)
- When user process needs to perform privileged action, must jump to privileged OS code, set CPU to high privilege, and then perform the action

Where is OS code located?

- OS is part of high **virtual address space of every process**
- Page table of process maps these kernel addresses to location of kernel code in memory
- Only one copy of OS code in memory, but mapped into virtual address space/page table of every process
- Process jumps to high virtual addresses to run OS code




Kernel mode execution

- When does a process go from user mode to kernel mode?
 - **System call**: user requests some privileged action from OS (e.g., read syscall)
 - **Program fault**: hardware raises a fault when user does some invalid action (e.g., privileged instruction in user mode, page fault)
 - **Interrupt**: I/O devices request attention (e.g., packet has arrived on NIC)
 - All these are called **traps** in general
- How is a trap handled?
 - Change CPU privilege level to high privilege/kernel mode
 - Jump to OS code that handles trap and run it
 - Return to user code after handling trap
 - Can choose to return to user code of another process too (if context switch)

What happens upon a trap?

- Trap handling begins by running a CPU instruction (`int n` in x86)
 - Invoked by system call code or triggered by hardware event
- What happens during execution of `int n`?
 - Change CPU privilege level
 - Lookup **Interrupt Descriptor Table (IDT)** with index “n” to get address of kernel code that handles this trap → set EIP to this value
 - Use **kernel stack** of process (in PCB) as the main stack → set ESP to this value
 - Start saving user context onto the (kernel) stack
- Next, kernel code to handle the trap runs
 - Save more user context, beyond what is saved by hardware instruction
 - System call processing, interrupt handling etc.
- Finally, kernel invokes `iret` (x86) instruction to return back to user mode
 - Reverses changes of `int n`

Trap handling



Interrupt Descriptor Table (IDT)

- Every interrupt has a number (IRQ)
 - E.g., different hardware devices get different numbers, system call gets a different number
- IDT has an entry for every interrupt number (IRQ)
 - IDT entry specifies values of EIP and a few such CPU registers
 - CPU uses this EIP to locate kernel interrupt handling code
- Pointer to IDT is stored in CPU register
 - Setting IDT pointer in CPU is a privileged operation, done by OS
 - Much like how setting CR3 to page table is a privileged operation

I/O subsystem: system calls, interrupts

- Processes use system calls to access I/O devices
 - Process P1 makes **system call**, goes into kernel mode
 - OS device driver initiates I/O request to device (e.g., disk, network card)
 - If the system call is **blocking** (i.e., cannot be completed right away), OS performs context switch to another process P2
 - When request completes, I/O device raises **interrupt**
 - P2 goes into kernel mode, handles interrupt, marks P1 as ready to run
 - P1 runs at a later time when scheduled by OS scheduler
- **DMA**: I/O devices perform Direct Memory Access to store I/O data in memory
 - When initiating I/O request, device driver provides (physical) address of memory buffer in which to store I/O data (e.g., disk block)
 - Device first stores data in DMA buffer before raising interrupt
 - Interrupt handler need not copy data from device memory to RAM

Summary of OS concepts

- The concept of a process, memory image, CPU context
 - CPU context is saved in PCB/kernel stack during user/kernel mode transitions of a process, as well as during context switch between processes
- Virtual memory, paging, address translation by MMU using page table
 - OS is mapped into the virtual address space of every process
 - Modern OSes use flat segments + paging
- User mode and kernel mode of a process, privileged instructions, CPU privilege levels
 - Privileged instructions in OS code run at highest CPU privilege level (ring 0)
- Interrupt/trap processing in kernel mode
 - Change privilege level, save user context, jump to kernel code, handle trap
- I/O handling